

Oído: Open-Vocabulary Speech Recognition on a \$5 Microcontroller without a Neural Accelerator

Daniel Varela Sánchez Guillermo Martín-Coello Juárez
Lokutor (Hashing Works, S.L.L.), Madrid, Spain
contact@lokutor.com

Abstract

Speech recognition on low-cost microcontrollers is, in practice, limited to keyword spotting or a fixed list of commands. We present Oído, an inference engine that runs open-vocabulary English speech recognition entirely on an ESP32-S3, a commodity microcontroller with a 240 MHz dual-core Xtensa LX7 CPU, 512 KB of on-chip SRAM, 8 MB of PSRAM and 16 MB of flash, and no neural accelerator. Oído executes a 13 M-parameter Conformer-CTC model with per-channel int8 weights, per-frame dynamic int8 activations and int8 relative-position attention, using the chip’s 128-bit SIMD instructions. Its execution schedule is designed around the memory bandwidth of this class of device (flash and PSRAM) rather than arithmetic. On LibriSpeech test-clean/test-other, the on-device arithmetic reaches 3.70%/8.23% word error rate (WER) with greedy decoding, within 0.12 points of the full-precision model, and 3.33%/7.15% with an on-chip 1.3 M-parameter recurrent language model. The same engine runs a Conformer-Transducer at 2.99%/6.67%, and an int4 variant that fits in 8.3 MB. On test-clean the int8 models make 2.3–2.8 $\times$ fewer errors than the vendor’s recognizer on the same chip, and they make fewer errors than Whisper tiny.en and Moonshine tiny running in full precision on a laptop, including under real recorded noise and reverberation. Measured on a physical ESP32-S3 board, the firmware transcribes at 1.97 $\times$ the duration of the audio on one core, with transcripts identical to those of an instruction-exact emulator, so the recognizer is accurate but not yet real time: the text of a 4 s command appears about 8.7 s after it ends. The two-core schedule, which emulation had suggested would approach real time, gives wrong transcripts on silicon and is disabled. A streaming variant, fine-tuned for chunked attention and trained with noise and reverberation, shows partial text while the speaker talks and, at this speed, delivers the final text 4.1–11.1 s after the end of a 3.6–10.6 s clip, against 7.2–22.4 s in utterance mode, at 4.86%/11.03% WER with the language model. A Spanish model for the same engine makes 34–60% fewer errors than multilingual Whisper tiny on four Spanish test sets. Code and models are released.

1 Introduction

Voice is the natural interface for speakers, appliances, toys, wearables and vehicles, but most of these devices are built around microcontrollers that cost a few dollars and cannot run the speech recognizers used on phones and servers. Today such products either stream audio to the cloud, paying for inference for the lifetime of every unit and exposing users’ audio, or they recognize a small, fixed set of wake words and commands on device [45]. Even the most capable vendor solution for the popular ESP32-S3 accepts only predefined command lists [14].

Accurate small speech models do exist. A 10–15 M-parameter Conformer [20, 30] reaches 2.7–3.7% WER on LibriSpeech test-clean, better than Whisper tiny [35], which has three times as many parameters. The obstacle to running such a model on a microcontroller is not its arithmetic. Flash is read through a small cache at tens of MB/s, and a model of this size occupies most of the flash and more than all of the on-chip SRAM. In addition, existing microcontroller runtimes [10] target static graphs with small activations, not an attention model that runs over whole utterances.

This paper shows that the gap can be closed on a commodity microcontroller with no neural accelerator. Our contributions:

- An int8 inference engine for Conformer speech recognition on the ESP32-S3. It loses 0.02 and 0.12 absolute WER points on LibriSpeech test-clean and test-other relative to full precision, without retraining (Section 5.1).
- A bandwidth-aware execution schedule: utterance-level row blocking, K-split accumulation and a batched front end. It reduces weight traffic from flash and PSRAM by 2.4 $\times$, and a set of kernel optimizations removes 9% of executed instructions (Section 5.3).
- On-chip decoding with a recurrent language model, and an RNN-T head whose joint network is evaluated for several frames at once while remaining exactly equivalent to greedy search (Section 5.4).
- A streaming mode: a Conformer fine-tuned for chunked attention and noise-robust, and an engine that caches int8 keys, values and convolution state per layer so that only the last partial chunk remains when the speaker stops.




Figure 1: Oído on the ESP32-S3: pipeline and memory placement. Weights are streamed from flash or PSRAM once per 64-frame block. The activations of the current block and the int32 accumulation tiles stay in on-chip SRAM.

We measure on silicon what it buys at the chip’s present speed (Sections 5.5 and 7.3).

- An evaluation on seven test sets, fourteen noise and reverberation conditions, voice-command latency and memory, against laptop-class models and the vendor’s on-chip recognizer (Section 7).

The engine, firmware, tools and models are open source.¹

2 Related Work

Small speech recognition models. Conformer [20] combines self-attention with convolution. Its 10.3 M-parameter variant reaches 2.7%/6.3% on LibriSpeech [32] with an RNN-T decoder. Squeezeformer [25] and Zipformer [43] improve the accuracy–compute trade-off. Whisper tiny [35] (39 M) and Moonshine tiny [23] (27 M) target laptops, phones and single-board computers. We use NVIDIA’s openly licensed Conformer-CTC Small [26, 30], which reaches 3.7%/8.1% greedy in full precision, and its transducer counterpart [29].

Quantization for speech recognition. Integer-only inference [22] is standard for edge deployment. For Conformers, Ding et al. [13] report native 4-bit quantization-aware training without accuracy loss on LibriSpeech, and Rybakov et al. [36] push to 2 bits (3.1%/7.0% at 10.5 MB, against 2.5%/6.1% in floating point). Zhen et al. [46] study sub-8-bit RNN-T and Conformer models. These works quantize with training and target accelerators or mobile CPUs. We show that post-training int8 with per-frame activation scales is essentially lossless for this architecture, and we use quantization-aware training only for int4.

Machine learning on microcontrollers. TensorFlow Lite Micro [10] and MLPerf Tiny [5] established the tooling and benchmarks for keyword spotting [45], visual wake words and anomaly detection, with models of tens to hundreds of kilobytes. For speech recognition, Arm has shown a wav2letter model on a Cortex-M55 with an Ethos-U55 neural accelerator [3, 4], and an int8 Conformer-S on the Ethos-U85 accelerator [40], simulated on Arm’s Corstone-320 platform and deployed by Alif on an Ensemble E8 board (Cortex-M55 with Ethos-U85). The latter reports 6.4% WER on test-clean for

the full-precision model. Espressif reports 8.5%/21.3% for its MultiNet7 recognizer on the ESP32-S3 [14], which in deployment decodes against a fixed command list. To our knowledge, no published microcontroller result, with or without a neural accelerator, reports a lower LibriSpeech WER.

3 Platform

The ESP32-S3 [16] integrates two Xtensa LX7 cores at 240 MHz with a single-precision FPU and the PIE 128-bit SIMD extension. The instructions we rely on are:

- `EE.VMULAS.S8.ACCX`, which multiplies two vectors of 16 int8 values and adds the 16 products into a 40-bit accumulator, optionally loading the next operand in the same instruction;
- `EE.VSMULAS.S8.QACC`, which multiplies a vector by a broadcast scalar into sixteen 20-bit lanes.

Our module (ESP32-S3-WROOM-1-N16R8)² has 512 KB of internal SRAM, 8 MB of octal PSRAM and 16 MB of quad-SPI flash. PSRAM and flash are accessed through a shared 64 KB data cache (configured size, taken from internal SRAM) with 64-byte lines. A cache miss to quad-SPI flash at 80 MHz transfers a line in roughly $2\ \mu\text{s}$, so sustained streaming from flash is on the order of 30–40 MB/s, and octal PSRAM is several times faster.

A 13 M-parameter model therefore has three placement options. At int8 it does not fit in SRAM at all. It fits in PSRAM only partially, because PSRAM also has to hold the working buffers. It fits in flash, but flash is slow to stream. The engine keeps the model in a flash partition mapped into the address space, copies the most frequently read matrices to PSRAM, and keeps per-block activations in SRAM.

4 Model

Oído runs the architecture of NVIDIA’s Conformer-CTC Small [30], summarized in Table 1. We re-implemented it in PyTorch from the released weights and matched the reference implementation: features agree to within 2×10^{-5} , and frame-level argmax labels agree exactly on the test utterances we checked.

¹<https://github.com/lokutor-ai/oido>; models at <https://huggingface.co/lokutor-ai>.

²US\$5.19 in single units and US\$3.41 in reels at LCSC (part C2913202), September 2026. The module includes the chip, PSRAM, flash and antenna.

Table 1: Model architecture (Conformer-CTC Small).

Features	80 log-mel, 25 ms window, 10 ms hop
Subsampling	$2 \times \text{conv } 3 \times 3$, stride 2, 176 ch $\rightarrow$ 25 Hz
Encoder	16 Conformer blocks, $d=176$, FF 704
Attention	4 heads, relative position [9]
Convolution	depthwise, kernel 31, BatchNorm
Output	CTC [19], 1024 SentencePiece [27] + blank
Parameters	13 M (int8: 14.0 MB; int4: 8.3 MB)

Processing is utterance-level. A voice activity detector (Section 5.6) delimits each utterance, and the encoder then runs over all of its frames with full attention. This matches the model’s training and avoids the accuracy cost of chunked streaming, at the price of latency: text appears after the speaker stops (Section 7.6).

5 Engine

5.1 Quantization

Every linear layer uses symmetric int8 weights with one scale per output channel, $W_q = \text{round}(W/s_w)$ with $s_w = \max_k |W_{nk}|/127$. Activations are quantized dynamically, with one scale per frame (row), computed at run time just before each matrix product:

$$x_q = \text{round}(x/s_x), \quad s_x = \frac{1}{127} \max_k |x_k|, \quad (1)$$

$$y_n = s_x s_{w,n} \sum_k x_{q,k} W_{q,nk} + b_n. \quad (2)$$

Accumulation is exact in int32. The rescaling and bias use fp32. Per-frame scales matter because activation magnitudes vary strongly between speech and silence frames; a single static scale per tensor would waste precision on quiet frames.

Self-attention is also integer. For head h , the content and position terms of Transformer-XL-style attention [9] are computed as two int8 dot products: $(q_i + u) \cdot k_j$ and $(q_i + v) \cdot p_{j-i}$. Queries (with the bias vectors u, v added), keys and projected positions each get per-(frame, head) scales. After the row maximum is subtracted, the softmax numerators are multiplied by the per-frame scales of the values $s_{v,j}$ and quantized relative to their row maximum. The value matrix is stored transposed in int8, so both the scores and the probability-value product are SIMD dot products, and the normalization by the softmax sum is applied once per output.

LayerNorm, SiLU, GLU gating and the depthwise convolution (with BatchNorm folded into its weights) run in fp32 on the FPU. The first convolution has a single input channel and also stays in fp32. The second is executed as an int8 GEMM over im2col patches with one scale per patch. Because the input to that layer is post-ReLU and non-negative, each patch’s maximum is the maximum over a 3×3 neighbourhood of per-column channel maxima, which the first convolution tracks as it writes its output. The patch quantization is therefore fused into the gather, with arithmetic identical to a separate pass.

5.2 Kernels

Matrix products use `EE.VMULAS.S8.ACCX`. For reduction lengths $K \leq 256$ (most layers, $K=176$) an *x-stationary* kernel keeps up to six 16-byte chunks of an activation row in vector registers while weight rows stream through the load slot of the multiply-accumulate instruction, which gives 16 MACs per instruction in steady state. For $K=704$ a row-streaming kernel is used, and for $K > 1024$ the K-split kernel of Section 5.3. The int4 variant unpacks nibbles with vector masks and uses `EE.VSMULAS.S8.QACC` as an outer product that produces 16 outputs per instruction; with int4 weights the 20-bit lanes cannot overflow for $K \leq 512$. The softmax uses a 1/64-step lookup table for e^{-x} , whose worst-case relative error (0.8%) is below the int8 probability step. Work can be split over both cores with a parallel-for over output channels, heads or frames; in the emulator, 93% of executed instructions fall in parallel regions. On silicon this mode currently gives wrong transcripts (Section 8.1), so every measurement in this paper uses one core.

5.3 Bandwidth-aware schedule

Each layer processes the utterance in blocks of 64 frames (2.56 s of audio). A weight row is loaded once per block and reused for all 64 frames while the block’s int8 activations stay in SRAM. Two cases break this pattern in a naive implementation.

- **Long reductions.** Layers with $K > 1024$ read activations from PSRAM. Row-streaming kernels must shrink the block so that the activations stay cache-resident, which multiplies weight reads. We instead split K into chunks of at most 512 and accumulate int32 partial sums for $16 \text{ outputs} \times 64 \text{ frames}$ in SRAM, so every weight byte is fetched once per block.
- **The front-end convolution.** It produces 20 patch rows per output frame. We batch three frames per GEMM.

With a simple traffic model (bytes of each weight matrix times reads per second of audio), these changes reduce weight streaming from about 18 MB/s to 7.5 MB/s, at a cost of 2.9% more executed instructions for the split accumulation.

Memory is budgeted per utterance:

- The positional-encoding table is quantized to int8 once and shared by all layers.
- The attention output buffer aliases the convolution-module buffer, since their lifetimes do not overlap.
- Front-end rows are quantized as they are produced.

The measured PSRAM working-set peak is 1.9 MB for a 15 s utterance and 2.4 MB for 20 s, the maximum utterance length. The remaining PSRAM holds the most frequently read weights, starting with the front end.

```

Input: joint encoder projections  $F_{1..T}$ ; prediction state  $g$ 
 $t \leftarrow 1$ 
while  $t \leq T$ :
   $Z \leftarrow \text{relu}(F_{t..t+B-1} + g)$ ;  $L \leftarrow \text{quant}(Z) W_{\text{out}}^\top$  (one
  pass over  $W_{\text{out}}$ )
   $i \leftarrow \text{first row of } L \text{ whose argmax} \neq \text{blank}$ 
  if none:  $t \leftarrow t + B$ ; continue
  emit argmax; update  $g$  with the LSTM; continue frame  $t+i$ 
  symbol by symbol (at most 5)
   $t \leftarrow t + i + 1$ 

```

Figure 2: Batched greedy transducer search. It is exactly equivalent to sequential greedy search, but reads the joint output matrix once per B blank frames instead of once per frame.

5.4 Decoding

CTC. Greedy decoding takes the per-frame argmax. With a language model, we run CTC prefix beam search [21] (beam 4, top-6 tokens per frame) with shallow fusion [24]: weight 0.3 and a token bonus of 0.5, tuned on dev-clean. The language model is a GRU [7] (embedding 192, hidden 384, 1.3 M parameters, int8 per-row weights) over the same 1024 tokens. We train two versions. The released one uses only the LibriSpeech language-model corpus [32] of public-domain books. A second one adds transcripts of our training corpora; it is not released because some of those transcripts carry share-alike terms. Language-model steps for all live beams are batched into one matrix product per frame.

Transducer. The same engine runs NVIDIA’s Conformer-Transducer Small [18, 29]. The encoder is architecturally identical; the head adds a 320-unit LSTM prediction network and a joint network with an output matrix of 1025×320 (328 KB). Greedy transducer search evaluates the joint network at least once per frame, which would stream that matrix from memory 25 times per second. Because the prediction state changes only when a token is emitted, and most frames emit only blank, we evaluate the joint network for up to $B=8$ frames at once with the current prediction state (Figure 2). Rows after the first frame that emits a token are discarded and recomputed with the updated state. Each row’s arithmetic is unchanged, so the output is identical to sequential greedy search. We verified this on the full LibriSpeech test sets: word errors, substitutions, deletions and insertions are identical.

int4 variant. For applications that need flash space, we fine-tuned the model with 4-bit quantization-aware training for 8,000 steps. Conformer-block linear layers use int4 per-channel weights. The front end, output head and positional projections stay int8, and activations and attention stay int8, as in the engine. Training used public corpora: Common Voice, VoxPopuli, LibriSpeech, MLS [34], AMI, VCTK [42], People’s Speech with transcripts regenerated by a larger model [31], and two small accent corpora from OpenSLR (Nigerian English, SLR70, and English accents of the British Isles, SLR83 [12]).

5.5 Streaming

Utterance mode must wait for the speaker to finish, then run the whole encoder. To remove that wait we fine-tune the same network for chunked processing, in the spirit of dynamic-chunk training [44] and cache-based streaming Conformers [28], with three changes that keep the pretrained weights usable. (i) *Chunked attention*: a frame attends to the frames of its own chunk and of the previous chunks within a left context; chunk sizes (8–32 frames of 40 ms) and left contexts (unlimited to 256 frames) are drawn per batch, and full context with probability 0.3, so one model runs in both modes. (ii) *Chunk-limited depthwise convolution*: the symmetric 31-tap kernel sees any past frame but nothing after the end of its chunk. Shifting the kernel to be causal instead needs far more retraining; this keeps the pretrained kernel. (iii) *Fixed feature normalization*: global per-mel statistics replace per-utterance statistics, which are unavailable before the utterance ends, with random gain augmentation. We also train with additive noise, music and babble from MUSAN [38] and simulated room impulse responses [37] (Section 6), which proved necessary: a first model fine-tuned on clean read speech lost much of the original’s robustness (Table 5).

The engine runs every C encoder frames through all layers with exactly the training masks. Per layer it caches the int8 keys, per-key scales and transposed values of the last $\lceil L/C \rceil$ chunks (L is the left context), the last 15 convolution inputs, and the positional projections, which depend only on a bounded range of relative offsets and are computed once. The caches have a fixed size (about 2 MB at $C=32$, $L=128$, kept in PSRAM) whatever the utterance length, and the transcript does not depend on how the audio is split into feeds (verified). When the speaker stops, only the last partial chunk is left.

This has a price in weight traffic. A chunk re-reads every weight, so traffic grows by $64/C$ relative to the 64-frame blocks of utterance mode (Section 5.3): $2\times$ at $C=32$ and $4\times$ at $C=16$. At the chip’s present speed this does not show: on the board the streaming engine runs at $1.96\times$ real time at $C=32$ on one core, slightly faster than the same model in utterance mode ($2.04\times$), because instruction issue rather than weight bandwidth limits a single core (Section 7.5). It may matter once both cores work. We default to $C=32$ (1.28 s).

5.6 System

Figure 1 shows the complete pipeline. The firmware captures 16 kHz audio from an I2S MEMS microphone into an 8 s PSRAM ring buffer. A segmenter applies a slow AGC and an energy voice activity detector against a tracked noise floor, which falls quickly and rises slowly, with 9/4.8 dB on/off hysteresis. An utterance ends after 0.8 s of silence or at 20 s. Each utterance is transcribed while capture continues, and the transcript is printed and optionally drawn on an I2C OLED display. The same C code builds for the host with `-ffp-contract=off` on both targets. The firmware on the board and in the emulator produce identical transcripts on all 27 clips we ran, and the activations of the two clips we

traced layer by layer are bit-identical. The host build links a laptop’s math library, which rounds the last bit of the log-mel features differently; that changes a word on some hard clips (the host agrees with the board on 22 of the 27, Section 7.5). All accuracy numbers in this paper come from the host build.

6 Evaluation Methodology

Data. We evaluate on the full LibriSpeech test-clean and test-other sets (2,620 and 2,939 utterances), and on held-out sets not used for training any model in this paper:

- Common Voice [2]: 1,500 utterances;
- VoxPopuli [41]: 1,200;
- AMI individual headset microphones [6]: 1,000;
- People’s Speech [17]: 1,000, with the corpus’s own, noisy references;
- Earnings-22 [11]: 4,179 segments.

NVIDIA’s training data includes Common Voice 7.0 and VoxPopuli [30], and Moonshine’s includes Common Voice, AMI and People’s Speech [23], so these models may have seen some of our held-out utterances in training. Laptop baselines are evaluated on up to 500 evenly spaced utterances per set. All systems share one segmentation and one text normalization (lowercasing, punctuation and filler removal), so some baseline numbers differ from those published by their authors, most visibly Moonshine tiny on Earnings-22.

Robustness. We mix 300 test-clean utterances with real noise recordings from DEMAND [39] (car, kitchen, living room, cafeteria) at 10 and 5 dB SNR, and with four-talker babble drawn from dev-clean. We also convolve them with simulated room impulse responses [37]: a $5\times 4\times 3$ m room, talker 2 m from the microphone, RT60 of 0.4 and 0.8 s. A far-field condition combines 0.6 s reverberation with kitchen noise at 10 dB.

Speed. We run the firmware on an ESP32-S3-WROOM-1-N16R8 module on a DevKitC-1-class board, with ESP-IDF 5.5.1, the CPU at 240 MHz, flash in QIO mode and octal PSRAM at 80 MHz, on one core (Section 8.1). Clips are stored in flash and transcribed one after another, each timed with the 64-bit microsecond timer; the real-time factor (RTF) is compute time divided by audio duration, so it includes the language-model beam search but not audio capture. (The 32-bit CPU cycle counter wraps every 17.9 s at 240 MHz and under-reports clips whose compute takes longer.) For the streaming engine we also feed each clip at the rate of speech and measure when the final text appears (Section 7.3). In addition we run the unmodified firmware in Espressif’s QEMU [15] with `-icount 0`, which advances virtual time by one nanosecond per executed instruction, summed over both emulated cores, so that the emulated cycle counter reads

Table 2: LibriSpeech WER (%). “Chip” rows use the on-device engine (host build, same int8 arithmetic). Laptop rows are full precision on 500-utterance subsets. [†]WER published by the vendor; size of the model files in ESP-SR.

System	Size	clean	other
Conformer-CTC S, fp32 (reference)	52 MB	3.68	8.11
Oído int8, greedy	14.0 MB	3.70	8.23
Oído int8 + GRU LM	15.3 MB	3.33	7.15
Oído int4 (QAT), greedy	8.3 MB	4.61	9.98
Oído int4 + GRU LM	9.6 MB	3.78	8.43
Oído Transducer int8, greedy	15.5 MB	2.99	6.67
MultiNet7 (ESP32-S3) [†]	2.7 MB	8.5	21.3
Whisper tiny.en (laptop)	39 M par.	6.26	15.93
Moonshine tiny (laptop)	27 M par.	5.00	12.13
Vosk small (laptop)	40 MB	9.95	21.61

exactly 1/25 of the executed instructions. We use it for exact instruction counts and as a functional reference: its transcripts equal the board’s, and dividing the board’s cycles by the emulator’s instructions gives the cycles per instruction of the real chip.

Baselines. Laptop baselines, all in full precision:

- Whisper tiny.en [35] (39 M parameters);
- Moonshine tiny [23] (27 M);
- Vosk small, a 40 MB Kaldi model for embedded Linux [1, 33].

For the ESP32-S3 we quote Espressif’s published MultiNet7 figures [14]; its evaluation protocol is not detailed.

7 Results

7.1 Accuracy and quantization

Table 2 shows that int8 post-training quantization with per-frame activation scales costs 0.02 (clean) and 0.12 (other) absolute WER points. The GRU language model reduces LibriSpeech errors by 10–13% relative. The transducer head brings the on-chip result to 2.99%/6.67%, in line with NVIDIA’s published full-precision 2.9%/6.6%. Plain int4 post-training quantization degraded accuracy substantially (to 14.1%/22.0% on the full test sets); after quantization-aware fine-tuning, the int4 model reaches 4.61%/9.98% greedy. On test-clean, every on-chip configuration makes 1.8–2.8× fewer errors than MultiNet7, and all int8 configurations make fewer errors than the laptop baselines.

Table 3 extends the comparison to other domains. The on-chip int8 models make fewer errors than all laptop baselines on Common Voice, VoxPopuli and AMI. On Earnings-22 they trail Whisper tiny.en slightly (20.8–22.4% vs 20.3%). Conversational meeting speech (AMI) and earnings calls remain hard for all models of this size.

Table 3: WER (%) on held-out sets. CV: Common Voice; Vox: VoxPopuli; PS: People’s Speech (noisy references); E22: Earnings-22. “Books”: LM trained on public-domain books only (released).

System	CV	Vox	AMI	PS	E22
<i>Oído on ESP32-S3 (on-chip engine)</i>					
int8, greedy	11.4	9.3	20.2	11.6	22.4
int8 + LM (books)	10.4	9.2	20.4	12.2	21.6
int8 + LM (books + transcripts)	10.1	9.0	19.0	11.0	20.8
int4 + LM (books)	11.8	11.1	20.2	10.7	21.1
Transducer int8	9.6	8.2	19.2	11.6	21.0
<i>Laptop baselines (fp32)</i>					
Whisper tiny.en	19.7	13.5	21.5	14.3	20.3
Moonshine tiny	17.5	13.4	24.2	8.1	32.4
Vosk small	27.1	20.6	34.1	12.9	38.0

The language model’s domain matters. The book-only model improves LibriSpeech, Common Voice and Earnings-22, leaves VoxPopuli unchanged, and slightly *hurts* AMI and People’s Speech relative to greedy decoding. The model that also saw speech transcripts helps on every set. For conversational products, the language model should be trained on conversational text.

7.2 Robustness

Table 4 reports WER under noise and reverberation. Car noise at 5 dB SNR costs the on-chip models less than 0.3 points, kitchen and living-room noise at 5 dB roughly 1–2 points, and cafeteria noise, which contains competing speech, 4.7–6.0 points. Averaged over the 14 conditions, the on-chip int8 models make 31–45% fewer errors than Whisper tiny.en and Moonshine tiny. The most difficult conditions are four-talker babble at 5 dB, strong reverberation and far-field use; at those distances, microphone arrays with beamforming would help more than a different recognizer.

7.3 Streaming

Table 5 compares the modes of one streaming-trained model with the utterance-mode model. Chunking costs accuracy, because the encoder cannot see what follows the chunk: with the language model, $C=32$ reaches 4.86%/11.03% on LibriSpeech test-clean/test-other against 3.33%/7.15% in utterance mode, and $C=16$ is worse still (6.14%/13.23%). The streaming model is still more accurate than the vendor’s recognizer (8.5%/21.3%) and than Whisper tiny.en (6.3%/15.9%) on a laptop. Its full-context mode (3.43%/7.78%) loses about 0.1–0.3 points to the original, the price of noise-robust fine-tuning on read speech.

Noise augmentation matters more than anything else we tried. Averaged over the 14 conditions of Table 4, the first streaming model, fine-tuned on clean speech with fixed normalization, scored 8.9% in full-context mode and 12.0% with $C=32$ (both with the language model), worse than the original model’s 7.5%. After adding MUSAN noise, music, babble

Table 4: WER (%) on 300 test-clean utterances under real noise (DEMAND), babble and simulated reverberation. T: Oído Transducer int8; C: Oído CTC int8 greedy; C+LM: with GRU LM; M: Moonshine tiny; W: Whisper tiny.en; V: Vosk small.

Condition	T	C	C+LM	M	W	V
Clean	2.5	3.5	3.3	4.3	5.6	9.1
Car, 10 dB	2.6	3.6	3.3	4.3	5.8	9.1
Car, 5 dB	2.7	3.6	3.2	4.4	5.9	9.3
Kitchen, 10 dB	3.0	4.0	3.6	5.1	6.6	10.7
Kitchen, 5 dB	3.3	4.7	4.2	6.1	7.5	12.4
Living room, 10 dB	3.2	4.4	4.0	5.8	6.7	11.5
Living room, 5 dB	4.2	5.7	4.8	7.0	7.9	15.1
Cafeteria, 10 dB	3.9	5.3	4.9	7.3	8.3	14.4
Cafeteria, 5 dB	7.2	9.5	8.6	12.2	12.3	24.7
Babble, 10 dB	8.4	10.9	9.7	14.1	14.5	36.7
Babble, 5 dB	25.9	27.6	24.7	37.3	35.9	70.9
Reverb, RT60 0.4 s	4.6	5.8	4.9	9.8	10.2	15.3
Reverb, RT60 0.8 s	12.4	16.0	14.1	28.6	22.5	33.6
Far-field kitchen	10.4	13.5	11.4	24.3	19.8	30.4
Mean (14 conditions)	6.7	8.4	7.5	12.2	12.1	21.7

Table 5: One streaming-trained model on the ESP32-S3 engine (int8, with the language model unless stated), against the original utterance-mode model. Robust.: mean WER over the 14 conditions of Table 4.

[†]Measured on the board, one core, six clips of 3.6–10.6 s: time from the last sample of the clip to the final text, with the streaming rows fed at the rate of speech and the utterance rows compute only. The 0.8 s end-of-speech wait comes on top.

Mode	clean	other	robust.	final text [†]
Utterance mode (original)	3.33	7.15	7.5	6.9–21.3 s
Streaming, $C=32$, greedy	6.27	13.06	–	–
Streaming, $C=32$	4.86	11.03	9.0	4.1–11.1 s
Streaming, $C=16$	6.14	13.23	–	–
Same model, full context	3.43	7.78	6.2	7.2–22.4 s

and simulated reverberation (8,000 further steps), the numbers fall to 6.2% and 9.0%: the streaming model is *more* robust than the original and, at $C=32$, about as robust as Whisper tiny.en (12.1%) is accurate.

Latency. Two numbers matter. When the chip has kept up with the audio, only the last partial chunk is left, and the final step takes 1.1 s on average on the board (0.8–1.6 s over 12 clips); in the emulator its instruction count is well described by $17\text{ M} + 5.5\text{ M} \times (\text{remaining encoder frames})$, summed over both cores. But the chip does not keep up: at $1.96 \times$ real time the backlog grows with the utterance. We therefore also fed six clips at the rate of speech, one 20 ms block when its last sample would have been spoken, and measured when the final text appears after the last sample: 4.1–11.1 s (mean 6.7 s) for clips of 3.6–10.6 s, against 7.2–22.4 s for the same clips in utterance mode (the 0.8 s end-of-speech wait comes on top of both). The delay grows with the length of the utterance, roughly as $(\text{RTF} - 1)T$ plus the final step. Streaming roughly halves the wait but does not give the 1.1–1.4 s we had estimated from

Table 6: Spanish WER (%) on 400-utterance subsets (Common Voice 17, MLS, VoxPopuli, FLEURS es_419).

System	CV	MLS	Vox	FLEURS
Oído + LM (ESP32-S3)	13.3	10.6	15.5	11.3
Oído greedy	20.3	14.2	20.0	16.9
Oído streaming $C=32$ + LM	15.2	12.1	16.3	12.8
Whisper tiny (laptop)	33.0	21.5	28.7	17.1

instruction counts; that needs a real-time factor below one.

7.4 Spanish

The engine is language-agnostic: a new language needs a new 1024-piece vocabulary, a re-initialized CTC head and fine-tuning of the whole network from the English weights. We trained a Spanish model on 2,492 h of Common Voice, VoxPopuli, MLS and FLEURS [8] (all CC0 or CC-BY), including 1,029 h of Common Voice’s unvalidated clips kept only where NVIDIA’s large Spanish Conformer [30] agrees with the reference within 15% WER, with the noise and reverberation augmentation of Section 5.5, and a Spanish GRU language model trained on public-domain books, Common Voice sentences and VoxPopuli transcripts with all development and test sentences removed. The model needs no engine change; the language model stores its own beam-search weights.

Table 6 compares it with multilingual Whisper tiny on the same utterances (400 evenly spaced per set, 333 for FLEURS after dropping references with digits; same normalization, numerals in Whisper’s output spelled out). With the language model it makes 34–60% fewer errors on all four sets, and streaming at $C=32$ costs about 2 points. The comparison favors our model on these domains, since we trained on the training splits of the corpora and Whisper is zero-shot, and we have no Spanish noise benchmark. On the complete test sets (60 h) the numbers are 13.75/10.90/15.68/11.32% with the language model. The language-model weights were chosen from a small grid on subsets of the test sets, on a flat optimum (all settings within 0.2 points of each other).

7.5 Speed and memory on the board

Table 7 gives the measured real-time factor. With one core, the int8 CTC model with the language model runs at $1.97\times$ real time on 27 clips of 3.6–10.6 s (1.91–2.02), the int4 model at 1.76, the transducer at 1.89 and the Spanish model at 2.06. The RTF barely depends on the clip, so compute time is proportional to the length of the audio. Figure 3 shows the executed instructions per second of audio in the emulator, measured on 39.7 s of test-other speech: 220 M across both cores, 46% of them in matrix products and 15% in attention. The board spends 461 M cycles per second of audio, 2.0 cycles per executed instruction; Table 8 shows that every stage costs between 1.3 and 2.4 cycles per instruction except the im2col gather (3.7). The time is thus spread evenly rather than concentrated in a memory-bound stage, and on one core the cost is set by

Table 7: Measured on the board, one core, 240 MHz, with the language model unless noted. RTF: compute time divided by audio duration (mean, and range over clips). “= host”: clips whose transcript equals the host build’s. The firmware’s transcripts equal the emulator’s on all 27 clips of the first row. Clips: 214.8 s of audio in the first row, 59–103 s in the others.

Model	clips	RTF	range	= host
CTC int8 (released)	27	1.97	1.91–2.02	22/27
CTC int4	10	1.76	1.73–1.83	8/10
Transducer int8, no LM	10	1.89	1.88–1.91	10/10
Stream. model, utterance	12	2.04	1.98–2.11	12/12
Stream. model, $C=32$	12	1.96	–	10/12
Spanish, utterance	12	2.06	1.98–2.13	9/12
Spanish, $C=32$	12	1.97	–	–

Table 8: Cycles on the board against instructions in the emulator, per second of audio (int8 CTC model, 27 clips on the board, 39.7 s of speech in the emulator).

Stage	Minstr/s	Mcyc/s	cyc/instr
log-mel	10	14	1.48
conv0	14	20	1.45
im2col	6	20	3.66
matrix products	103	216	2.09
layernorm	7	15	2.07
activation quant	15	32	2.16
SiLU / GLU	11	25	2.35
q/k/v int8	6	13	2.04
pos. enc.	3	6	2.03
attention	34	44	1.30
depthwise conv	10	15	1.40
head + beam search with LM	21	40	1.90

how many instructions the kernels issue, not by waiting for weights. Beam search with the language model takes 8.7% of the cycles. If both cores were perfectly parallel, the 55% of instructions on the critical path would give an RTF of about 1.1; this is a projection, and the one build we could time with both cores ran at 0.78 times the one-core time (Section 8.1).

Table 9 lists the engine optimizations and their effect on executed instructions. The bandwidth-aware tiling deliberately increases instructions, by 2.9%, to cut weight traffic $2.4\times$. Its benefit is invisible to an instruction-counting emulator, and we did not measure it on silicon (that needs a build without the tiling).

Flash holds 14.0 MB (int8), 8.3 MB (int4) or 15.5 MB (transducer), plus 1.3 MB for the optional language model. With int4, a 16 MB module leaves a 6 MB partition for application code. At the maximum utterance length the measured PSRAM working set is 2.4 MB, and at least 125 KB of internal SRAM stays free for the rest of the system.

7.6 Latency

Table 10 gives the time from the last sample of a short voice command to its text, in each mode. In utterance mode it is




Figure 3: Executed instructions per second of audio, by stage, for the int8 CTC model (both cores, exact counts from instruction-counting emulation).

Table 9: Engine optimizations: executed instructions for 39.7 s of audio (int8 CTC, greedy), cumulative. [†]Adds instructions but cuts weight traffic from flash from 18 to 7.5 MB/s.

Step	G instr.	Δ
Baseline int8 engine	9.82	—
+ merged attention score pass	9.63	−1.9%
+ x-stationary GEMM for $K \leq 256$	9.39	−2.5%
+ K-split, batched front end [†]	9.66	+2.9%
+ int8 positional table, buffer aliasing	9.58	−0.8%
+ fused im2col quantization	9.29	−3.0%
+ softmax lookup table	8.99	−3.2%
+ two-frame depthwise convolution	8.95	−0.4%
Total		−8.9%

the compute time, which scales with the length of the clip at the measured RTF; the 0.8 s end-of-speech wait comes on top. The streaming rows come from the run fed at the rate of speech (Section 7.3). At present, a 3.7 s command takes about 7 s in utterance mode and 4 s in streaming mode, and neither is conversational.

8 Discussion and Limitations

Where the accuracy comes from. The acoustic models are NVIDIA’s, trained on NeMo ASRSET, a composite of public English corpora that NVIDIA’s model cards put at about 16,000 hours [30]. Our contribution is to run them on a microcontroller without an accelerator and with almost no loss: per-frame dynamic quantization keeps int8 essentially lossless without retraining, and whole-utterance, block-wise execution lets a 14 MB model run from flash. We designed the schedule around memory traffic, amortizing every weight fetch over 64 frames. The board shows, however, that once the model fits and runs, the remaining cost on one core is instruction issue: about two cycles for each instruction in every stage, with the matrix kernels spending most of their instructions on loads, accumulator reads and loop overhead rather than on multiply-accumulates. Faster kernels, and the second core, are where the next factor of two lies.

Table 10: Time from the last sample of a 3.6–3.8 s command to its text, measured on the board (one core). [†]Computed from the measured RTF (this model was not run on a clip of that length).

Profile	RTF	text after clip
CTC int8 + LM, utterance mode	1.97	6.9–7.4 s
CTC int4 + LM, utterance mode	1.76	6.2–6.7 s
Transducer int8, utterance mode	1.89	7.0 s [†]
Streaming model, utterance mode	2.04	7.2–7.6 s
Streaming model, streaming $C=32$	1.96	4.1–4.2 s

8.1 The two-core mode

In the emulator the two-core schedule is correct, and its instruction counts suggested a real-time factor of 0.7–0.95. On the board, at 240 MHz, it gives wrong transcripts and sometimes crashes. We traced it with a hash of every layer’s output on one 3.8 s clip (each configuration was run three times). One core is exactly reproducible: all 166 hashes are identical in every run and identical to the emulator’s. With two cores the activations are corrupted at random places, always first in layer 1 or 2 and almost always in the feed-forward block of layer 1, and in roughly one run in eight the firmware crashes with an `IllegalInstruction` exception and a corrupted backtrace, which suggests corrupted instruction or data fetches rather than a logic error. Two cores at 160 MHz are correct (166 of 166 hashes identical, in all three runs), but no faster than one core at 240 MHz.

Parallelizing every part of the model except the matrix products is correct, and parallelizing only the matrix products fails. Among them, the $K=176$ products and the $K=3520$ front-end products are correct on two cores, and the $K=704$ products (the second linear layer of each feed-forward block), which move the most weight data per cycle, corrupt. Neither PSRAM speed (80 to 40 MHz), data-cache line size (64 to 32 B), the location of the weights (flash or PSRAM), FreeRTOS semaphores (a spin-wait barrier behaves the same), interrupts (masked while a job runs), the location or sharing of the activation tile, instruction fetch (kernel in IRAM, or a separate copy per core), nor the split itself (by rows instead of by output channels) removes the errors. We do not know the cause. A hardware margin problem at 240 MHz (supply droop or timing) fits the 160 MHz result, the random values and the crashes, but a second engine of ours (a text-to-speech engine) running a synthetic two-core stress of the same shape on the same board at 240 MHz ($K=704$, 64 rows, a shared 45 KB tile in internal SRAM, an 88/88 split, 3×90 s, plus other shapes, about 21 minutes in all, every call checked against a one-core reference) saw no mismatch and no crash, and its real firmware passed a 600-iteration soak. A plain supply problem on this board therefore looks unlikely; something specific to our kernel, to how the two cores are started and joined, or to streaming weights from external memory in lock-step is more likely, and we have not found it. Even a correct two-core mode would not reach real time on this evidence: the incorrect spin-wait build ran at 0.78 times the one-core time, not the 0.55 that the instruction counts promised, because the cores

contend for memory.

8.2 Limitations

- **Not real time yet.** On one core the recognizer runs at about twice real time ($1.97\times$): the text of a 4 s command appears about 8.7 s after it ends in utterance mode (including the 0.8 s end-of-speech wait) and about 5 s after it in streaming mode. The two-core mode that would roughly halve this is not correct on silicon (Section 8.1).
- **Streaming costs accuracy and, at this speed, buys a factor of two in latency, not more.** Chunked processing loses about 1.5 points on test-clean against utterance mode, and the final text still arrives roughly $(RTF - 1)T$ after an utterance of length T .
- **Small board study.** We measured one board. The RTF figures cover 27 clips (CTC int8) and 10–12 clips for the other models, all of at most 10.6 s; the 2.4 MB PSRAM peak for a 20 s utterance comes from the emulator. Accuracy numbers come from the host build; the board’s transcripts equal the emulator’s, and the host differs from them on some hard clips (22 of 27 identical), most visibly for Spanish (the board scores 9.3% against 7.3% for the host build on 12 clips, 205 words).
- **Two languages.** English and Spanish; other languages need a new vocabulary and fine-tuning.
- **Transducer license.** The transducer weights are distributed under NVIDIA’s NGC terms and are not redistributed with our release.
- **Comparisons.** The laptop baselines use 500-utterance subsets, and MultiNet7 numbers are vendor-reported.
- **Framework comparison.** We did not compare against TensorFlow Lite Micro or ESP-DL deployments of the same model.

9 Conclusion

Oído runs open-vocabulary English speech recognition entirely on a \$5 microcontroller with no neural accelerator, at 3.70%/8.23% LibriSpeech WER (2.99%/6.67% with a transducer head). That is within 0.12 points of full precision and more accurate than full-precision laptop models of two to three times the size. On the board it runs at about twice real time on one core, so today it suits short commands where several seconds of delay are acceptable, not live transcription or conversation. A streaming variant roughly halves the wait after the speaker stops, at a cost in accuracy, and the two-core mode, which is not yet correct on silicon, and faster kernels are the open routes to real time. We release the engine, firmware, models and the board measurements so that others can build on them and check them.

Acknowledgments

We thank NVIDIA for releasing the NeMo Conformer models under open terms, and the creators of the datasets used for evaluation. Parts of the engine code, the experiment scripts and this manuscript were drafted with the help of an AI assistant (Claude, Anthropic). The authors directed the work, checked the results and the text, and take full responsibility for the content.

References

- [1] Alpha Cephei. Vosk speech recognition toolkit: vosk-model-small-en-us-0.15. <https://alphacephei.com/vosk/models>, 2020. Accessed September 2026.
- [2] Rosana Ardila, Megan Branson, Kelly Davis, Michael Kohler, Josh Meyer, Michael Henretty, Reuben Morais, Lindsay Saunders, Francis Tyers, and Gregor Weber. Common voice: A massively-multilingual speech corpus. In *Proc. LREC*, 2020.
- [3] Arm. ML embedded evaluation kit: Automatic speech recognition use case. <https://gitlab.arm.com/artificial-intelligence/ethos-u/ml-embedded-evaluation-kit>, 2026. Accessed September 2026.
- [4] Arm. Arm ML-Zoo: Wav2letter. https://github.com/Arm-Examples/ML-zoo/tree/master/models/speech_recognition/wav2letter, 2026. Accessed September 2026.
- [5] Colby Banbury, Vijay Janapa Reddi, Peter Torelli, Jeremy Holleman, Nat Jeffries, Csaba Kiraly, Pietro Montino, David Kanter, Sebastian Ahmed, Danilo Pau, et al. MLPerf Tiny benchmark. In *Proc. NeurIPS Datasets and Benchmarks Track*, 2021.
- [6] Jean Carletta et al. The AMI meeting corpus: A pre-announcement. In *Machine Learning for Multimodal Interaction (MLMI 2005)*, volume 3869 of *Lecture Notes in Computer Science*, pages 28–39. Springer, 2006. doi: 10.1007/11677482_3.
- [7] Kyunghyun Cho, Bart van Merriënboer, Caglar Gulcehre, Dzmitry Bahdanau, Fethi Bougares, Holger Schwenk, and Yoshua Bengio. Learning phrase representations using RNN encoder-decoder for statistical machine translation. In *Proc. EMNLP*, 2014.
- [8] Alexis Conneau, Min Ma, Simran Khanuja, Yu Zhang, Vera Axelrod, Siddharth Dalmia, Jason Riesa, Clara Rivera, and Ankur Bapna. FLEURS: Few-shot learning evaluation of universal representations of speech. In *Proc. IEEE SLT*, 2022.
- [9] Zihang Dai, Zhilin Yang, Yiming Yang, Jaime Carbonell, Quoc V. Le, and Ruslan Salakhutdinov. Transformer-XL: Attentive language models beyond a fixed-length context. In *Proc. ACL*, 2019.
- [10] Robert David, Jared Duke, Advait Jain, Vijay Janapa Reddi, Nat Jeffries, Jian Li, Nick Kreeger, Ian Nappier, Meghna Nataraj, Shlomi Regev, Rocky Rhodes, Tiezheng Wang, and Pete Warden. TensorFlow Lite Micro: Embedded machine learning for TinyML systems. In *Proc. MLSys*, 2021.

- [11] Miguel Del Rio, Peter Ha, Quinten McNamara, Corey Miller, and Shipra Chandra. Earnings-22: A practical benchmark for accents in the wild. *arXiv preprint arXiv:2203.15591*, 2022.
- [12] Isin Demirsahin, Oddur Kjartansson, Alexander Gutkin, and Clara Rivera. Open-source multi-speaker corpora of the English accents in the British isles. In *Proc. LREC*, pages 6532–6541, 2020.
- [13] Shaojin Ding, Phoenix Meadowlark, Yanzhang He, Lukasz Lew, Shivani Agrawal, and Oleg Rybakov. 4-bit conformer with native quantization aware training for speech recognition. In *Proc. Interspeech*, 2022.
- [14] Espressif Systems. ESP-SR benchmark (ESP32-S3). <https://docs.espressif.com/projects/esp-sr/en/latest/esp32s3/benchmark/README.html>, 2026. Accessed September 2026.
- [15] Espressif Systems. Espressif QEMU. <https://github.com/espressif/qemu>, 2026. Release esp_develop_9.2.2_20260417. Accessed September 2026.
- [16] Espressif Systems. ESP32-S3 technical reference manual. https://documentation.espressif.com/esp32-s3_technical_reference_manual_en.pdf, 2026. Version 1.8. Accessed September 2026.
- [17] Daniel Galvez, Greg Diamos, Juan Ciro, Juan Felipe Cerón, Keith Achorn, Anjali Gopi, David Kanter, Maximilian Lam, Mark Mazumder, and Vijay Janapa Reddi. The people’s speech: A large-scale diverse English speech recognition dataset for commercial usage. In *Proc. NeurIPS Datasets and Benchmarks Track*, 2021.
- [18] Alex Graves. Sequence transduction with recurrent neural networks. *arXiv preprint arXiv:1211.3711*, 2012.
- [19] Alex Graves, Santiago Fernández, Faustino Gomez, and Jürgen Schmidhuber. Connectionist temporal classification: Labelling unsegmented sequence data with recurrent neural networks. In *Proc. ICML*, 2006.
- [20] Anmol Gulati, James Qin, Chung-Cheng Chiu, Niki Parmar, Yu Zhang, Jiahui Yu, Wei Han, Shibo Wang, Zhengdong Zhang, Yonghui Wu, and Ruoming Pang. Conformer: Convolution-augmented transformer for speech recognition. In *Proc. Interspeech*, 2020.
- [21] Awni Y. Hannun, Andrew L. Maas, Daniel Jurafsky, and Andrew Y. Ng. First-pass large vocabulary continuous speech recognition using bi-directional recurrent DNNs. *arXiv preprint arXiv:1408.2873*, 2014.
- [22] Benoit Jacob, Skirmantas Kligys, Bo Chen, Menglong Zhu, Matthew Tang, Andrew Howard, Hartwig Adam, and Dmitry Kalenichenko. Quantization and training of neural networks for efficient integer-arithmetic-only inference. In *Proc. CVPR*, 2018.
- [23] Nat Jeffries, Evan King, Manjunath Kudlur, Guy Nicholson, James Wang, and Pete Warden. Moonshine: Speech recognition for live transcription and voice commands. *arXiv preprint arXiv:2410.15608*, 2024.
- [24] Anjuli Kannan, Yonghui Wu, Patrick Nguyen, Tara N. Sainath, Zhifeng Chen, and Rohit Prabhavalkar. An analysis of incorporating an external language model into a sequence-to-sequence model. In *Proc. ICASSP*, 2018.
- [25] Sehoon Kim, Amir Gholami, Albert Shaw, Nicholas Lee, Kartikeya Mangalam, Jitendra Malik, Michael W. Mahoney, and Kurt Keutzer. Squeezeformer: An efficient transformer for automatic speech recognition. In *Proc. NeurIPS*, 2022.
- [26] Oleksii Kuchaiev, Jason Li, Huyen Nguyen, Oleksii Hrinchuk, Ryan Leary, Boris Ginsburg, Samuel Krizan, Stanislav Beliaev, Vitaly Lavrukhin, Jack Cook, Patrice Castonguay, Mariya Popova, Jocelyn Huang, and Jonathan M. Cohen. NeMo: a toolkit for building AI applications using neural modules. *arXiv preprint arXiv:1909.09577*, 2019.
- [27] Taku Kudo and John Richardson. SentencePiece: A simple and language independent subword tokenizer and detokenizer for neural text processing. In *Proc. EMNLP: System Demonstrations*, 2018.
- [28] Vahid Noroozi, Somshubra Majumdar, Ankur Kumar, Jagadeesh Balam, and Boris Ginsburg. Stateful Conformer with cache-based inference for streaming automatic speech recognition. In *Proc. ICASSP*, 2024.
- [29] NVIDIA. STT en conformer-transducer small, version 1.6.0. https://catalog.ngc.nvidia.com/orgs/nvidia/teams/nemo/models/stt_en_conformer_transducer_small, 2022. Accessed September 2026.
- [30] NVIDIA. NVIDIA Conformer-CTC small (en-US). https://huggingface.co/nvidia/stt_en_conformer_ctc_small, 2023. CC-BY-4.0. Accessed September 2026.
- [31] NVIDIA. parakeet-tdt-0.6b-v2. <https://huggingface.co/nvidia/parakeet-tdt-0.6b-v2>, 2025. CC-BY-4.0. Accessed September 2026.
- [32] Vassil Panayotov, Guoguo Chen, Daniel Povey, and Sanjeev Khudanpur. Librispeech: An ASR corpus based on public domain audio books. In *Proc. ICASSP*, 2015.
- [33] Daniel Povey, Arnab Ghoshal, Gilles Boulianne, Lukas Burget, Ondrej Glembek, Nagendra Goel, Mirko Hannemann, Petr Motlicek, Yanmin Qian, Petr Schwarz, Jan Silovsky, Georg Stemmer, and Karel Vesely. The Kaldi speech recognition toolkit. In *Proc. IEEE ASRU*, 2011.
- [34] Vineel Pratap, Qiantong Xu, Anuroop Sriram, Gabriel Synnaeve, and Ronan Collobert. MLS: A large-scale multilingual dataset for speech research. In *Proc. Interspeech*, 2020.
- [35] Alec Radford, Jong Wook Kim, Tao Xu, Greg Brockman, Christine McLeavey, and Ilya Sutskever. Robust speech recognition via large-scale weak supervision. In *Proc. ICML*, 2023.
- [36] Oleg Rybakov, Phoenix Meadowlark, Shaojin Ding, David Qiu, Jian Li, David Rim, and Yanzhang He. 2-bit conformer quantization for automatic speech recognition. In *Proc. Interspeech*, 2023.
- [37] Robin Scheibler, Eric Bezzam, and Ivan Dokmanić. Pyroomacoustics: A python package for audio room simulation and array processing algorithms. In *Proc. ICASSP*, pages 351–355, 2018. doi: 10.1109/ICASSP.2018.8461310.

- [38] David Snyder, Guoguo Chen, and Daniel Povey. MU-SAN: A music, speech, and noise corpus. *arXiv preprint arXiv:1510.08484*, 2015.
- [39] Joachim Thiemann, Nobutaka Ito, and Emmanuel Vincent. The diverse environments multi-channel acoustic noise database (DEMAND): A database of multichannel environmental noise recordings. In *Proc. Meetings on Acoustics*, volume 19, page 035081, 2013. doi: 10.1121/1.4799597.
- [40] Alexey Vilkin, Alex Tawse, Kshitij Sisodia, and George Gekov. End-to-end int8 Conformer on Arm: Training, quantization, and deployment on Ethos-U85. <https://developer.arm.com/community/arm-community-blogs/b/internet-of-things-blog/posts/end-to-end-int8-conformer-on-arm-training-quantization-and-deployment-on-ethos-u85>, 2026. Arm Community blog, 21 January 2026. Accessed September 2026.
- [41] Changhan Wang, Morgane Riviere, Ann Lee, Anne Wu, Chaitanya Talnikar, Daniel Haziza, Mary Williamson, Juan Pino, and Emmanuel Dupoux. VoxPopuli: A large-scale multilingual speech corpus for representation learning, semi-supervised learning and interpretation. In *Proc. ACL*, 2021.
- [42] Junichi Yamagishi, Christophe Veaux, and Kirsten MacDonald. CSTR VCTK corpus: English multi-speaker corpus for CSTR voice cloning toolkit (version 0.92). University of Edinburgh, The Centre for Speech Technology Research (CSTR), 2019.
- [43] Zengwei Yao, Liyong Guo, Xiaoyu Yang, Wei Kang, Fangjun Kuang, Yifan Yang, Zengrui Jin, Long Lin, and Daniel Povey. Zipformer: A faster and better encoder for automatic speech recognition. In *Proc. ICLR*, 2024.
- [44] Binbin Zhang, Di Wu, Zhuoyuan Yao, Xiong Wang, Fan Yu, Chao Yang, Liyong Guo, Yaguang Hu, Lei Xie, and Xin Lei. Unified streaming and non-streaming two-pass end-to-end model for speech recognition. *arXiv preprint arXiv:2012.05481*, 2020.
- [45] Yundong Zhang, Naveen Suda, Liangzhen Lai, and Vikas Chandra. Hello edge: Keyword spotting on microcontrollers. *arXiv preprint arXiv:1711.07128*, 2017.
- [46] Kai Zhen, Martin Radfar, Hieu Duy Nguyen, Grant P. Strimel, Nathan Susanj, and Athanasios Mouchtaris. Sub-8-bit quantization for on-device speech recognition: A regularization-free approach. In *Proc. IEEE SLT*, 2022.